
COMP2017 & COMP9017 Week 1 Tutorial

Introduction to Linux

Introduction to Linux

What is Linux

Linux is a family of operating systems that derive from the original AT&T Unix developed in the 1970's at Bell Labs. As an operating system it manages the computer's hardware and software and provides a common interface for applications to use as they manipulate elements of the system.

Unix follows a philosophy that revolves around designing small, well defined programs that integrate well together. This can be summarised as follows:

- Write programs that do one thing and do it well
- Write programs to work together
- Write programs to handle text streams, because that is the universal interface

Everything in Unix is either a file or a process. A file is any collection of data, such as a series of configuration strings, an image, or compiled source code, while a process is a program that is currently being executed. Unix then revolves mostly around passing and parsing streams of data between different processes and files.

The Kernel and the Shell

The Unix kernel is a master control program that handles low level tasks such as starting and stopping programs, managing the file system and peripheral devices and passing instructions to processors or graphics cards. The kernel is complimented by a series of small programs that perform common tasks such as listing the files in a directory, changing directories and executing compiled code. Other tools can then build off these functions to provide search functions, document readers and image processors such as desktop environments.

Booting into Linux

If your lab machine is currently running windows, shut it down and restart it. At start-up should see a boot menu allowing you to switch between Windows and Red-Hat Linux. Using the arrow keys, select Linux and press enter to boot.

Upon reaching the login screen press **Ctrl + Alt + F2**, your screen should turn black and you should be shown a login prompt. Login using your uni-key and password.

Question 1: Getting the Tutorial Sheet

At this point you should only have access to a black screen and a prompt. In order to get your tutorial sheet you're going to have to use git: **git clone** <https://github.com/comp2017/tutorial1.git>. To read it from the command line we're going to want to use one of these tools that converts pdf to readable text; **pdftotext**. This will output a text file scraped from the pdf of the tutorial sheet.

In order to use this file efficiently you are going to need to read a restricted range of lines from the file, in particular using the **cat** and the **sed** commands.

- Use the git command to clone the repository containing the tutorial sheet to your computer.
- Use the pdftotext command to strip the text from the pdf and put it into another file
- The **sed -n 1,3p <file>** command should print lines 1 to 3 of the text file that you specify, this can be changed for any range of lines that you require. This will be invaluable for you throughout this tutorial. If you did not specify a file name when you entered the pdftotext then it will default from <file>.pdf to <file>.txt.

Terminal Basics

One of the most useful linux commands is **man**. This brings up the **manual** pages for a given command and shows options and example usages if you are unsure of the syntax. To **quit** the man page, press **q**. As an example, enter the command **man ls** to see the manual pages for the **list** command. This command shows the names of all the files and folders in the current directory. You should see in the man page the option **ls -a** which shows all files.

Quit out of the man page and try both **ls** and **ls -a**. You should notice that the second command lists quite a few more files than the first, all of them prefixed with a full stop.

mkdir is the **make directory** command. As the command suggests it creates a new directory with whatever name you specify. For example **mkdir comp2017** creates a new directory in the current folder called 'comp2017'. Directories can be removed with the **rmdir** command.

Next we want to be able to navigate between directories. For this we have the **change directory** command **cd**. To move to the comp2017 directory you would simply **cd comp2017**.

Tab Completion

Typing out the full names of each folder every time is somewhat redundant and to work around this Linux features tab completion. Pressing tab will automatically complete your current command or path if it can be uniquely determined. If it cannot be uniquely determined, pressing tab twice will list all possible completions.

From the comp2017 directory type **cd ../D** and then press tab twice. You should see the following output.

```
Desktop/ Documents/ Downloads/
```

Changing this to **cd ../Doc** and then pressing tab once should complete the command to **cd ../Documents**. Once you've entered this command use tab completion to navigate back to your comp2017 directory.

If your screen is looking full you can use the **clear** command or *Ctrl + L* to clear the screen.

Question 2: Terminal Navigation

Using **cd ..** navigate upwards from your current directory to the topmost or 'root' directory, pay attention to the path that you took and then try to navigate back to your home directory. If you get lost **cd ~** will automatically take you back.

Files

As stated above, in Unix everything is either a file or a process. Files can be read, written to and executed. To see what permissions a file has use the **-l** option for the **ls** command.

```
-rw-r--r-- 1 root root 387 Dec 10 10:17 Makefile
```

Here this Makefile has read and write permissions for the user, and read permissions for group and anyone. File permissions can be changed using the **chmod** command.

There are a number of commands that can be used to read and write to files, the most basic of which are **cat** and **touch**. The cat command **concatenates** files and prints the output, allowing the easy reading of plaintext files while touch changes the timestamps of a file, if the file doesn't exist then it creates it. **less** also reads text files, but places the user in an environment that is virtually identical to the one for **man**. It can be escaped using *q* in the same manner.

Existing files can be copied using **cp** moved **mv**, and removed **rm**. Directories are removed using **rmdir**. Though there are a few flags that can be passed to the rm command to grant it the same utility. Hidden files are prefixed with a **.** and can be viewed using the **-a** flag with the ls command. These hidden files are also called 'dotfiles' and tend to be configuration files for various programs.

To add text to a file we can use the **echo** command. Echo prints whatever input it is given, for example

```
echo "Hello World!"
```

We can redirect the output of any Linux command to any other Linux command, or to a file.

As per the Linux philosophy, text can be directed between processes and to files. In this case the **>** and **»** commands write and append to files respectively.

```
echo "Hello World!" > hello.txt
```

As all inputs and outputs from processes are text based, the output from one command can be ‘piped’ as the input to another command.

```
cat my_file.txt | less
```

The above command takes the output of cat and displays it using the ‘less’ command, which works in a similar manner to the **man** command you saw earlier.

Vim and Emacs

To edit a file from the command line, there exist two common text editors. These are **vim** and **emacs** and you should learn to use at least one of them.

Vim is smaller, and somewhat less featured while Emacs rivals operating systems in size but is more bloated.

Before you open and use either of them, it is incredibly useful to know in advance how to close them again. In vim you will be looking to press the escape button and then to enter the ‘:q’ command, or the ‘:q!’ command if you’re really stuck.

In Emacs you’re going to need to press ‘Ctrl + X Ctrl + C’, then either save or clear the buffer and answer ‘yes’ if prompted whether you’d like to exit anyway. In the emacs prompt control is abbreviated as a capital C.

To use either to edit a file simply enter **vim my_file** or **emacs my_file**.

In vim press the ‘i’ key to enter ‘insert’ mode, and type what text you need to. Then press escape to return to the ‘command mode’ from where you can save with ‘:w’ and quit with ‘:q’. You can also combine these commands to ‘:wq’ to save and close the file and vim.

In emacs you can immediately type your text into the ‘buffer’ which will be periodically auto-saved to the file. You can also force it to save using ‘C-x C-s’ or ‘Ctrl + x Ctrl + s’, and then quit.

This is a very brief introduction to these two editors and you should practice using them and learning other commands and shortcuts for each in your own time.

Directories, Devices and other Files

Folders or Directories are also files, opening a directory using vim or emacs will show some metadata about how the system interprets the directory along with a listing of the path to the directory and the files within it. The path to the current working directory of your terminal can be displayed using the **pwd** command.

Not all files are physically stored on disc. For example the current time is stored entirely in memory as a file found at */proc/uptime*. Use cat to view the contents of that file. What happens when you cat the file a second time?

Devices are also treated as files, with **lspci** showing devices connected on a PCI port, **lsusb** showing devices on a USB port and **lsblk** showing block devices. These commands also show the path to the files that represent the connection to these devices. By reading and writing to these files it is possible to read and write to and from the device itself.

Another useful command is **du** which will show the current disk usage of different files and folders, allowing for easy cleaning when the disk starts to fill.

Question 3: Drives and Network Shares

Have a look at the */etc/fstab* file (don't try to change it!). This file tells Linux what devices are to be mounted automatically, what the UUID of the device is so that it can be uniquely identified and what format the device takes.

- How many drives are currently mounted on your computer? Where is your home directory actually located?
- What other mounted drives look interesting? Rummage around in the other drive and see what you can find.
- What do you think the purpose of this second mounted network share is? Why has the system been designed this way?

Question 4: Who Administrates the Administrators?

The */home* directory contains not just your own folders and files, but also those of every other user on the system. As there are several thousand such users, it is ill advised to attempt to print them all to your terminal.

- Redirect the output of the list of folders in the */home* directory to a text file in your own home directory using.
- Using the **grep** command, search the file for the user names of all the administrators on the SIT network.
- Try doing the same using the **cat** command to pipe the file to the **grep** command.
- Using **grep** search your *.bash_history* file for the each time you used **grep**.

Question 5: Global Regular Expression Print

You may notice that the thing you are searching for using **grep** is displayed in red. Have a look in */etc/profile.d/colorgrep.sh*.

- What is really happening when you type the **grep** command?

- Create or edit your *.bashrc* file to create a new alias similar to the ones you saw in *colorgrep.sh* for a *grep_no_colour* command. The “*–color=NONE*” option will be helpful here.
- Run **source .bashrc** to run the commands in your *.bashrc* profile and try your new aliased command, have you de-coloured *grep*?

Question 6: Mounting and Unmounting

If you have a USB, plug it into your lab machine and using *lsblk* and the *mount* command, mount and access the contents of the USB. The *umount* command can be used to safely unmount a mounted device. Be careful not to unmount the drive containing the operating system! It is a common convention to mount temporary devices to */tmp*.

- Where is the *mount* command located?
- What are the permissions on the *mount* command? What looks odd about these permissions?

The UNIX Filesystem

In the root directory you can find a number of different folders with specific functions.

- *home* Contains the home directory for each user. Your home directory contains your Documents, Downloads and other directories. When logged in the system variable *\$HOME* is associated with your home directory. Try **cd \$HOME** to see it being used.
- *bin* contains binary executable files that execute most of the processes on the system. Having a look in here you should see a number of core unix commands such as *ls*, *touch*, *cat* and *su*.
- *boot* is often a separate boot partition rather than just a directory, you can check this using **lsblk**. This partition manages how the kernel is loaded when the computer boots. You should not modify anything in this folder unless you are sure that you know what you are doing, a mistake can render your operating system non-functional.
- *lib* contains common shared libraries and kernel modules, such as firmware and cryptographic functions.
- *dev* is the device folder. This contains files that communicate between the operating system and any connected devices. If you plug a usb into your Linux system you will find and be able to mount it from here.
- *root* The home folder for the root user, otherwise identical to a regular user's home folder.

Processes

A process is simply a currently executing program. Each process is assigned a unique process ID (PID) which can be used as a handle to start and stop it. A currently running process in the foreground can also be stopped using *Ctrl + C*.

The **ps** command displays a list of currently running processes under the current terminal process. The **-au** option for **ps** will list all processes on the system, including those of other users. For a more interactive display the **top** or **htop** commands can be used, these will also show the load on the CPU, the amount of RAM in use, the number of threads currently running and other useful diagnostic information.

Processes are spawned by passing a file with execution permissions to the kernel. The terminal you are currently using is itself a process and the commands you type are interpreted by the process as paths to executable files that are then run.

Your particular Linux distribution comes with a number of executable programs that perform common tasks such as changing directories, making and reading files. Some more advanced processes can be run in the background so that you can continue to use your terminal as they run. In order to background a process simply end the line containing the command' with an amperstand **&**. A backgrounded process can be foregrounded again using the **fg** command.

Question 7: Starting, Killing and Backgrounding Processes

The **ping** command is a small network utility that sends packets to a target IP address and waits for a response. It is quite useful for checking if another computer is connected to the network. You can check your own IP address on each of the networks that your computer is connected to using the **ip addr** command.

- Ping your own computer using the **ping** command, if you are unsure how to use **ping** use the **man** command.
- Stop the **ping** process using *Ctrl + C*.
- Start **ping** as a background process.
- Use the **kill** command with the process ID from the **ping** command to kill the **ping** process without foregrounding it.

Shell Scripts

A shell script is a file with execution permission that contains a series of bash commands. When executed it simply runs each line of the bash script as a command line command. This is a convenient method of re-using bash commands, however the question arises as to how the kernel knows to interpret this executable file as a bash script rather than, for example, a ruby script.

For this linux uses a string prefixed with a hash bang at the start of the file to indicate what command the file is to be run with. In the case of a shell script this is **#!/bin/bash**. The program loader then runs

the file with the command located along the specified path, in this case `/bin/bash` is the binary file that runs the bash terminal.

Question 8: Shell Script Basics

Write a shell script that prints your username by looking at the path to your home directory. You may find `~` or `$HOME` and the `sed` command helpful here.

Shell Scripts Continued

Bash is a Turing complete language in and of itself and comes with a number of the usual programmatic conveniences such as variables, switches and loops. It also has a number of different syntactic conventions and useful pre-defined constants such as `$USER` and `$HOME`. Try echoing those constants in the command line to see their current values.

Variables are also indicated by a `$` prefix when they are called, but not when they are initialised. As spacing is a strict syntax structure in bash, spaces around variable initialisations cannot be ignored else an error will be thrown. As Linux is built around text streams, all variables can be considered to be strings.

```
foo="bar"
echo foo
echo $foo
```

The dollar sign or grave quotes can also be used to encapsulate another bash command and pass it to a variable. As the output of all commands is a text stream, there are no issues surrounding typing.

```
foo=$(ls)
echo $foo
bar=`ls`
echo $bar
```

Bash switches are bookended by `if` and `fi` with the general syntactic structure following ‘`if [condition]; then else fi`’. Double equals signs are not used, and the terms being compared should be encapsulated in double quotes. Boolean expressions for use in switch statements and loops are encapsulated in square brackets.

```
if [ "$foo" = "$bar" ];
then
    echo sesquipedalian
else
    echo loquacious
fi
```


Bash loops operate over collections in a manner quite similar to Python. They follow a syntactic structure of

```
for file_var in `ls` ;  
do  
    echo file: $file_var  
done
```

Bash also makes use of while and until loops: while [condition]; do done, and: until [condition]; do done. To emulate traditional loops over a range of integers the **seq** command can be used to provide a collection over some range to loop over.

```
for i in seq 1 10;  
do  
    echo \"$i  
done
```

Alternatively, the **let** command can be used to change the value of a variable

```
foo=0  
while [ \"$foo -lt 10 ];  
do  
    echo \"$foo  
    let foo=foo+2  
done
```

Once again the lack of spacing when assigning a value to a variable is essential. Bash does support ++, += and related operators.

Bash scripts also accepts arguments, these are automatically set to variables \$1 for the first argument, \$2 for the second and so on.

Question 9: Your own Text Scroller

Write a bash script that tracks your place for the purposes of viewing a selected number of lines of a text file. Don't forget the hashbang!

- The script should read from a configuration file stored in an appropriately named directory in .config indicating what file it is to be reading from and the last read line.
- The script should have a flag that changes the file that it is reading
- When invoked using the -n <number> flag the script should read the next number of lines specified.
- When invoked using the -s flag the script should re-print the same

- When invoked using the `-p <number>` flag, the script should print the previous number of lines specified.
- When invoked without a flag, the script should set the new target file to whatever file the user has specified as an argument and reset the last read line to 0.
- Alias your script in `.bashrc` with the absolute path to the script to use the command without needing to replicate the path from your current directory to where the script is stored

If you had sufficient permissions, instead of aliasing your script you could instead provide a sim link from the `/bin` directory. For the sake of posterity, you can check the word count of your script using the `wc` command, and the line count using the `wc -l` command.

Desktop Environments

The window manager itself in Linux is a process that is managed by the kernel and provides a desktop environment with extravagances like mouse support, icons and windows. The window manager can be run using the **startx** command.

Find what the `startx` command does and what files might you edit to modify the behavior of the window manager. How does this relate to `xinit`?

Run the `startx` command to return to the safe and familiar land of desktop environments.

You should by this point see that the Unix operating system is simply a series of files that communicate through processes to run just about everything the operating system does. The flexibility afforded by this scheme means that modifying any of these files allows a user to tailor the system to their own ends and it has been squeezed onto just about every device containing an x86 processor that has been made.